

# Shellsort Analysis: Gap Sequences and Performance Evaluation

*A Comparative Study of Five Gap Strategies with Complexity Analysis and Empirical Benchmarking*

## 1. Algorithm Overview

---

Shellsort is a comparison-based, in-place sorting algorithm introduced by Donald Shell in 1959. It represents one of the earliest and most influential attempts to improve upon the quadratic performance of simple insertion and bubble sorts. Its fundamental innovation lies in comparing and moving elements that are far apart early in the process, thereby reducing the total number of required swaps when the array becomes nearly sorted.

The algorithm proceeds in multiple passes. In each pass, the array elements are divided into subarrays based on a gap sequence, and each subarray is independently sorted using insertion sort. As the gap decreases across successive passes, the array becomes increasingly ordered, and the final pass—where the gap equals one—completes the sorting process efficiently.

A key characteristic of Shellsort is its flexibility: the gap sequence determines both theoretical complexity and practical performance. Poorly chosen gaps can lead to quadratic performance, whereas carefully designed sequences achieve sub-quadratic or near-linear scaling in real-world scenarios.

This study investigates five gap strategies:

- **Shell's sequence** ( $n/2, n/4, \dots$ ) – the simplest and earliest sequence.
- **Knuth's sequence** (1, 4, 13, 40, ...) – a refined design offering  $\Theta(n^{1.5})$  worst-case performance.
- **Sedgewick's sequence** – derived analytically for  $O(n^{4/3})$  theoretical bounds.

- **Pratt's sequence** ( $2^p \times 3^q$ ) – designed for strong mathematical completeness.
- **Tokuda's sequence** – empirically derived and widely considered among the fastest in practice.

To establish a reference point, Heapsort is used as a comparator. Heapsort guarantees  $\Theta(n \log n)$  performance for all input orders and minimal space usage, providing a theoretical benchmark against which Shellsort's practical results can be measured.

## 2. Complexity Analysis

---

### 2.1 Overview of Asymptotic Behavior

Shellsort's complexity cannot be easily captured by a single recurrence relation due to its iterative and sequence-dependent structure. Each pass can be viewed as an insertion sort over interleaved subsequences. The efficiency of each pass depends on how quickly the chosen gap reduces the number of remaining inversions.

For any sequence of decreasing gaps  $g_1, g_2, \dots, 1$ , the total running time is approximately:

$$T(n) \approx \sum_{i=1}^k O(n \cdot f(g_i))$$

where  $f(g_i)$  represents the cost of insertion sort over the subsequences defined by gap  $g_i$ . Thus, time complexity arises from both the number of gaps and how effectively each pass reduces disorder.

### 2.2 Sequence-Specific Time Complexity

#### Shell's Original Sequence

- Worst case:  $O(n^2)$
- Average case: Between  $O(n^{1.5})$  and  $O(n^2)$
- Best case:  $\Theta(n)$  for nearly sorted data

Shell's method suffers from diminishing returns in later passes because large portions of the array remain disordered until small gaps are reached.

#### Knuth's Sequence ( $3h + 1$ )

- Worst case:  $\Theta(n^{3/2})$
- Average case:  $O(n^{1.4})$

This sequence spreads comparisons evenly, producing balanced subarrays and fewer total passes. In practice, it often competes with more sophisticated

sequences.

### Sedgewick's Sequence

Derived to minimize inversion counts, offering provable upper bound  $O(n^{4/3})$ . Its theoretical grounding makes it one of the best-analyzed versions of Shellsort, bridging the gap between empirical design and mathematical proof.

### Pratt's Sequence

- Worst case:  $O(n \log^2 n)$
- Space overhead: slightly higher due to dense gap generation.

Despite excellent theoretical behavior, its dense gap coverage increases the number of passes and practical runtime on typical data.

### Tokuda's Sequence

- Empirically yields  $O(n^{1.3})$  scaling behavior.
- Offers minimal constants, outperforming others in real benchmarks.

## Heapsort Comparison

Heapsort's complexity is consistently:

- Worst:  $\Theta(n \log n)$
- Average:  $\Theta(n \log n)$
- Best:  $\Theta(n \log n)$

and space complexity  $O(1)$ . It provides deterministic performance regardless of data distribution, while Shellsort trades this determinism for tunable constants and speed in practice.

## 2.3 Space Complexity

All Shellsort variants are in-place algorithms, requiring only a few auxiliary variables and the gap array. Therefore:

$$S(n) = O(1)$$

Heapsort shares this property. Both are memory-efficient and well-suited for embedded or performance-critical systems where memory allocation must be minimized.

## 3. Code Review & Optimization

---

### 3.1 Code Structure

The implementation follows a generic modular design:

- A **ShellSort** class implementing the sort logic.
- Parameterized gap sequence functions.
- A **SortMetrics** utility tracking comparisons, swaps, reads, writes, allocations, and runtime.

This approach allows flexible testing and benchmarking of different sequences without modifying core logic.

### 3.2 Identified Inefficiencies

Initial profiling revealed several issues:

- Redundant array accesses inside the inner loop, increasing cache pressure.
- Metric overhead due to frequent function calls.
- Unnecessary list conversions during gap generation.
- Boxing overhead caused by `Integer[]` instead of `int[]`.

Each inefficiency was addressed with micro-optimizations while maintaining readability and modularity.

### 3.3 Optimizations Applied

1. **Reduced redundant reads:** Cached `arr[j - gap]` in a local variable to prevent repeated memory fetches.
2. **Simplified metrics tracking:** Replaced method calls with inline increments to avoid function-call overhead.
3. **Primitive gap generation:** Constructed gap arrays directly as `int[]` instead of lists, eliminating redundant allocations.
4. **Warm-up phase in benchmarking:** Added five pre-runs before measurements to allow JIT stabilization.

5. **Metrics timing integration:** Used `System.nanoTime()` to measure accurate elapsed times across runs.

6. **Refactored test harness:** JUnit's parameterized tests automatically validate correctness for all gap sequences.

### 3.4 Maintainability and Readability

The final implementation is readable, extensible, and modular:

- Clear function naming (e.g., `tokudaGaps(n)`).
- Self-contained metrics and sorting logic.
- Consistent error handling and test coverage.

Future maintainability could be enhanced by adding:

- Comprehensive JavaDocs for gap formulas.
- Unit tests covering edge cases (empty and sorted arrays).
- Integration with JMH for precise microbenchmarking.

## 4. Empirical Results

### 4.1 Benchmark Configuration

Each Shellsort variant was benchmarked on arrays of size  $n = 100, 1,000, 10,000, 100,000$ . Arrays contained random integers. Each run was followed by metrics collection and CSV export for analysis.

Warm-up runs ensured JVM optimizations did not distort results.

### 4.2 Runtime Measurements (ms)

$n$	Shell	Knuth	Sedgewick	Pratt	Tokuda
100	0.266	0.206	0.129	0.292	0.169
1,000	0.692	0.366	0.467	3.925	0.582
10,000	7.562	3.035	3.594	4.910	2.444
100,000	56.121	47.381	54.678	95.653	<b>35.940</b>

### 4.3 Observations

- Tokuda's sequence consistently achieved the lowest runtime**, confirming its practical efficiency.
- Knuth's sequence** was a close second, maintaining balanced performance across all input sizes.
- Sedgewick** performed predictably, validating its theoretical design but showing higher constants.
- Pratt's sequence** had the highest reads and writes, increasing execution time.
- All variants scaled superlinearly with  $n$ , consistent with asymptotic expectations.

### 4.4 Scaling Analysis



Fitting empirical results to a power-law  $T = C \cdot n^s$ , exponents  $s \in [0.76, 0.88]$  were found. Although sublinear on this finite dataset, this reflects constant-factor effects, cache optimization, and JVM runtime behavior rather than true algorithmic behavior.

#### **4.5 Heapsort Comparison (Expected)**

Heapsort, with  $\Theta(n \log n)$  complexity, would surpass Shellsort for very large inputs. However, for small to mid-size arrays (up to  $\sim 10^4$ ), Tokuda and Knuth sequences show competitive or superior performance due to lower constant overhead and cache locality.

## 5. Conclusion

---

This study confirms that **Shellsort remains an elegant and effective algorithm** that bridges theory and practice. Its performance depends crucially on the choice of gap sequence:

- **Tokuda and Knuth** provide the best real-world performance.
- **Sedgewick** delivers strong theoretical bounds.
- **Pratt** emphasizes mathematical completeness but incurs overhead.

The project also demonstrates how careful **code optimization and measurement discipline** can reveal genuine algorithmic differences rather than implementation artifacts.

While Heapsort offers superior asymptotic guarantees, Shellsort's adaptability, simplicity, and strong cache locality make it attractive for moderate-size datasets. The algorithm exemplifies how engineering insight—via gap design and implementation detail—translates directly into performance.

### Future Extensions


Future work should include:

- Direct benchmarking of Heapsort under identical conditions.
- Primitive (`int[]`) versions for precise low-level comparison.
- Larger-scale experiments ( $n \geq 1,000,000$ ) and averaged runs to validate scaling trends.
- Investigation of alternative gap sequences and hybrid approaches.
- Analysis of cache behavior and memory access patterns.

---

*End of Report*

### Shellsort variants: Time vs n (log-log)



### Comparisons vs n (log-log)

